

SpMM vs GEMM Implementation and Performance Analysis

1 Question 1: Dense Matrix-Matrix Multiplication (GEMM)

My GEMM implementation employs several optimizations for high-performance matrix multiplication:

1.1 1. Tiling with Triple Blocking

- **Implementation:** Used three tiling parameters:
 - `TILE_ROW = 32` : Block size for rows of matrix A
 - `TILE_COL = 64` : Block size for columns of matrix B
 - `TILE_INNER = 128` : Block size for inner dimension
- **Purpose:** Optimizes cache utilization by ensuring working sets fit in L1/L2 cache

1.2 2. Memory Management

- **Aligned Memory Allocation:** `aligned_alloc(32, bytes)` for all matrices
- **Buffer Copying:** Uses temporary buffers `buffer_A` and `buffer_B` to improve locality
- **Pre-zeroing:** `std::fill_n(matrix_C, rows * cols, 0.0f)` for result matrix

1.3 3. AVX2 SIMD Vectorization

- **Implementation:**
 - 256-bit registers processing 8 floats per instruction
 - Key AVX2 instructions:

```
1 __m256 sum = _mm256_loadu_ps(&matrix_C[i_idx * cols + j_idx]);
2 __m256 a_val = _mm256_set1_ps(row_A[k_idx]);
3 __m256 b_vec = _mm256_loadu_ps(&buffer_B[k_idx * TILE_COL + (j_idx -
    j_start)]);
4 sum = _mm256_fmadd_ps(a_val, b_vec, sum);
```

- **Purpose:** 8x throughput for floating-point operations

1.4 4. Thread-Level Parallelism

- **Implementation:**

```
1 int num_threads = omp_get_num_procs();
2 omp_set_num_threads(num_threads);
3 #pragma omp parallel
```

- **Dynamic Scheduling:** `#pragma omp for schedule(dynamic, 1)` for load balancing

1.5 5. Scalar Fallback

- **Implementation:** Scalar code paths for edge cases (tiles not multiple of 8):

```

1 for (int j_offset = 0; j_idx + j_offset < j_end; ++j_offset) {
2     float* result_cell = &matrix_C[i_idx * cols + j_idx + j_offset];
3     float cell_sum = *result_cell;
4
5     for (int k_idx = 0; k_idx < k_end - k_start; ++k_idx) {
6         cell_sum += row_A[k_idx] * buffer_B[k_idx * TILE_COL + (j_idx -
7             j_start) + j_offset];
8     }
9     *result_cell = cell_sum;
10 }

```

2 Question 2: Sparse Matrix-Matrix Multiplication (SpMM)

My SpMM implementation efficiently handles sparse matrices using several advanced techniques:

2.1 1. Sparse Accumulator (SPA) Data Structure

- **Implementation:** Hash-based accumulator with optimized collision handling:

```

1 class SparseAccumulator {
2     int* indices;
3     double* values;
4     bool* occupied;
5     size_t capacity, size, table_size, mask;
6     inline void add(int col, double val) { ... }
7     std::vector<std::pair<int, double>> get_sorted_entries() { ... }
8 };

```

- **Purpose:** Efficiently accumulates sparse products without expensive searches

2.2 2. Sparsity-Aware Memory Management

- **Implementation:**

- Output arrays dynamically allocated based on estimated sparsity:

```

1 int estimate_row_nnz(const int* A_row_ptr, int row,
2                     const double* A_values, const int* A_col_ind,
3                     const int* B_row_ptr, int k, int n) {
4     double A_density = static_cast<double>(A_nnz) / k;
5     double B_avg_density = B_total_nnz / (k * n);
6     double expected_density = A_density * B_avg_density * 1.5 * k;
7 }

```

- Working only with non-zero elements

2.3 3. Prefetching for Irregular Memory Access

- **Implementation:** Strategic prefetching to hide memory latency:

```

1 if (jA + 1 < A_row_ptr[i+1]) {
2     __builtin_prefetch(&A_col_ind[jA + 1], 0, 3);
3     __builtin_prefetch(&A_values[jA + 1], 0, 3);
4 }
5
6 __builtin_prefetch(&B_row_ptr[col_A], 0, 3);

```

- **Purpose:** Mitigates cache misses from irregular memory access patterns

2.4 4. Thresholding for Numerical Stability

- **Implementation:** Filters out tiny values to prevent excessive fill-in:

```
1 if (std::abs(product) >= 1e-14) {
2     spa.add(col_B, product);
3 }
```

and

```
1 C_temp[i] = spa.get_sorted_entries(1e-14);
```

- **Purpose:** Maintains sparsity by dropping near-zero results

2.5 5. Adaptive Parallelization

- **Implementation:**

```
1 int chunk_size = std::max(1, std::min(64, m / (omp_get_max_threads() * 2)))
;
2 #pragma omp parallel for schedule(dynamic, chunk_size)
```

- **Purpose:** Balances load while minimizing thread management overhead

3 Question 3: Performance Comparison

3.1 Theoretical Analysis

| | GEMM | SpMM |
|----------------------|---|--|
| Storage | $O(n^2)$ for $n \times n$ matrices | $O(nnz)$ where nnz = non-zeros |
| Memory Access | Regular, predictable pattern | Irregular, pointer-chasing pattern |
| Operations | Always $O(n^3)$ for $n \times n$ matrices | $O(nnz(A) \times avg_nnz_per_row(B))$ |
| When to Use | Dense matrices ($\geq 10\%$ zeros) | Sparse matrices ($\geq 90\%$ zeros) |

Table 1: Theoretical comparison between GEMM and SpMM

3.2 Benchmarking Results from Our Tests

| Matrix used | Matrix size | Non zero entries | spmm | gemm |
|-------------------------------|-------------|------------------|-------|--------|
| dwt_512 (suit_sparse) | 512×512 | 3,502 | 10ms | 21ms |
| delaunay_n10 (suit_sparse) | 1024×1024 | 6,112 | 36ms | 53ms |
| rdb2048 (suit_sparse) (graph) | 2048×2048 | 12,032 | 74ms | 195ms |
| Facebook (SNAP dataset) | 4039×4039 | 0.54% sparsity | 341ms | 1047ms |

Table 2: Performance comparison between SpMM and GEMM implementations

4 Conclusion

My implementation of SpMM significantly outperforms GEMM for sparse matrices, with performance gains directly proportional to matrix sparsity. The key techniques contributing to this efficiency are:

1. The SparseAccumulator hash table for efficient accumulation
2. Strategic memory prefetching
3. Adaptive thread parallelism with intelligent chunking
4. Sparsity-preserving thresholds for numerical values

These optimizations make SpMM the clear choice for sparse matrices, especially in domains like network analysis, scientific computing, and machine learning with sparse features.